

# Hypothesis Testing

## Answering Questions about Data

### Assignment Contents

- [Calculate CI Using Python](#)
- [The T-distribution](#)
- [Confidence Intervals: Difference of Means](#)

### Overview

#### EXPECTED TIME 2 HRS

This assignment will review the calculation of confidence intervals and p-values in Python. The examples from lecture will be reproduced using Python functions and packages. Eventually you will be asked to calculate confidence intervals and p-values on your own.

### Activities in this Assignment

- Calculate confidence intervals
  - For means
  - For differences of means
  - Using t- and z- distributions
- Calculate p-values

**NOTE: The z-multipliers MUST be calculated with the stats package as demonstrated below. e.g. for a 95% confidence interval, 1.959963... must be used rather than just 1.96**

A quick reminder about a couple of useful "numpy" functions.

```
import numpy as np

# Calculations of:
print("Taking the mean:",
      np.mean([1,2,3,4,5,6,7]))
print("Finding Standard Deviation:",
      np.std([1,2,3,4,5,6,7]))
print("Finding Square-root:",
      np.sqrt(1738))
```

### Calculate CI Using Python

Lets go through the example given in Lecture 7-1 around the 8 minute mark:

**Observations:**  $n = 56$

**Sample Standard Deviation:**  $s = 25$

**Population Standard Deviation:** UNKNOWN

Likelihood that  $\bar{X}$  (the sample mean) is within  $\pm 5$  of  $\mu$ ?

$$P(-5 \leq \bar{X} - \mu \leq 5) = P(-5 \leq \text{Error} \leq 5)$$
$$P\left(\frac{-5}{\sigma/\sqrt{n}} \leq Z \leq \frac{5}{\sigma/\sqrt{n}}\right)$$

Using sample standard deviation as estimate of  $\sigma$ :

$$P\left(\frac{-5}{25/\sqrt{56}} \leq Z \leq \frac{5}{25/\sqrt{56}}\right)$$
$$P\left(\frac{-5}{3.34...} \leq Z \leq \frac{5}{3.34...}\right)$$
$$P(-1.49666 \leq Z \leq 1.49666)$$

Because of the symmetry of the normal distribution, this is equal to:

$$= 1 - 2 * P(Z \geq 1.49)$$
$$= 1 - 2 * (1 - P(Z \leq 1.49))$$
$$= 1 - 2 * P(Z \leq 1.49)$$

(Rounding used to match lecture)

Note above how  $P(Z \geq 1.49)$  is changed to  $P(Z \leq -1.49)$ . The reason for this is  $P(Z \leq y)$  maybe easily calculated for any value of  $y$  with the ".cdf()" function in the stats package.

Show n below :

```
import scipy.stats as stats
1 - 2 * stats.norm.cdf(-1.49)
```

Continuing the example, creating the confidence interval.

The interval is:

$$\bar{X} \pm Z_{\alpha/2} \cdot \frac{\sigma}{\sqrt{n}} \quad \bar{X} \pm Z_{\alpha/2} \cdot \frac{\sigma}{\sqrt{n}}$$

To calculate  $Z_{\alpha/2}$  for a 95% confidence interval we can use the ".interval()" function or the ".ppf()" function. Both of which were covered more extensively in earlier assignments.

```
print(stats.norm.interval(.95))
```

```
print(stats.norm.ppf( 1- ((1 - .95)/2)))
print(stats.norm.ppf( (1 - .95)/2))
```

Calculating the confidence interval:

```
alpha = .95
interval_end = 1-((1-alpha)/2)
print(interval_end)
z_mult = stats.norm.ppf(interval_end)
sd = 25
x_bar = 135
n = 56

# Using our calculated values from above and python to create CI
print("\nConventional Calculation:\n",
      (x_bar - z_mult*(sd/np.sqrt(n)), x_bar + z_mult*(sd/np.sqrt(n)) ))

# Calculating CI using the .interval function
# (FOR THOSE INTERESTED IN SEEING MORE OF THE STATS PACKAGE)
print("\nCalculation with .interval():\n", # "loc" is the mean of the distribution, "scale" is the sd
      stats.norm.interval(alpha = .95, loc = x_bar, scale= sd/ np.sqrt(n)))
```

## Question 1

```
### GRADED

### Given point estimate (x_bar), confidence level(Z[alpha/2]), n, and sample standard deviation,
### calculate a confidence interval.

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### e.g. In the above example, the correct answer would be:

# lower = 128.45222
# upper = 141.54778

### Answers will be tested to three decimal places

### Calculate a 95% confidence interval where the sample mean of 92 observations was 130 with a
### sample standard deviation of 12

### YOUR ANSWER BELOW

x_bar = 130
n = 92
alpha = .95
sd = 12
interval_end = 1-((1-alpha)/2)

z_mult = stats.norm.ppf(interval_end)

lower = x_bar - z_mult*(sd/np.sqrt(n))

upper = x_bar + z_mult*(sd/np.sqrt(n))
```

## Question 2

```
### GRADED

### Given point estimate (x_bar), confidence level(Z[alpha/2]), n, and sample standard deviation,
### calculate a confidence interval.

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places

### Calculate a 90% confidence interval where the sample mean of 22 observations was 150 with a
### sample standard deviation of 40

### YOUR ANSWER BELOW

x_bar = 150
n = 22
alpha = .90
sd = 40

interval_end = 1-((1-alpha)/2)

z_mult = stats.norm.ppf(interval_end)

lower = x_bar - z_mult*(sd/np.sqrt(n))

upper = x_bar + z_mult*(sd/np.sqrt(n))
```

## Question 3

```
### GRADED

### Given point estimate (x_bar), confidence level(Z[alpha/2]), n, and sample standard deviation,
### calculate a confidence interval.

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places
```

```

### Calculate a 99% confidence interval where the sample mean of 2000 observations was 140 with a
### sample standard deviation of 40

### YOUR ANSWER BELOW

x_bar = 140
n = 2000
alpha = .99
sd = 40

interval_end = 1-((1-alpha)/2)

z_mult = stats.norm.ppf(interval_end)

lower = x_bar - z_mult*(sd/np.sqrt(n))

upper = x_bar + z_mult*(sd/np.sqrt(n))

```

## Question 4

```

### GRADED

### Given point estimate (x_bar), confidence level(Z[alpha/2]), n, and sample standard deviation,
### calculate a confidence interval.

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places

### Calculate a 95% confidence interval where the sample mean of 40 observations was 120 with a
### sample standard deviation of 9

### YOUR ANSWER BELOW

x_bar = 120
n = 40
alpha = .95
sd = 9

interval_end = 1-((1-alpha)/2)

z_mult = stats.norm.ppf(interval_end)

lower = x_bar - z_mult*(sd/np.sqrt(n))

upper = x_bar + z_mult*(sd/np.sqrt(n))

```

## The T-distribution

When  $\sigma$  is unknown, and "n" is small: Lecture 7-3

The "stats" package has a library for the t distribution.

The "t" library functions similarly to the "norm" library, except that degrees of freedom must be specified. Remember, degrees of freedom (df) in these cases is  $n - 1$ . Thus 21 observations would yield  $df = 20$

The below show s how to calculate the values from the t-distribution discussed and looked up in the t-table in lecture 7-3.

```

print("normal distributions; 2.5% in tails; interval",
      stats.norm.interval(.95))

print("\nt-distribution; 2.5% in tails; df = 20; interval",
      stats.t.interval(.95, df = 20))

print("\nt-distribution; 2.5% in tails; df = 5; interval",
      stats.t.interval(.95, df = 5))

print("\nt-distribution; 2.5% in tails; df = 5; ppf",
      stats.t.ppf(.975, df = 5))

```

### Example:

Given a set of observations, and a confidence level of 95%, calculate the confidence interval with a t-distribution:

```

observations = [121, 110, 126, 112, 129, 118, 126, 127, 126, 111, 127, 113, 126, 115, 114, 116]

n = len(observations) # find "n" -- the number of observations
x_bar = np.mean(observations) # find "x_bar"-- the sample mean
sd = np.std(observations) # find the sample standard deviation
alpha = .95
t_mult = stats.t.interval(alpha, df = n-1)[1]

print("Sample Mean: ", x_bar)
print("Observations (n): ", n)
print("Sample sd: ", sd)
print("t-multiplier: ", t_mult)

print("\nConfidence Interval: ", x_bar, "+/-", round(t_mult * (sd / np.sqrt(n)),4))

```

Thus, given the above 16 observations, we can calculate the 95% confidence interval from a t-distribution; spanning from ~116.3 to ~123.3

## Question 5

```
### GRADED

### Calculate the 95% confidence interval (with a t-distribution)
### of the data stored in the "observations" variable below

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places

### YOUR ANSWER BELOW

observations = [104, 148, 109, 104, 108, 120, 134, 129, 140, 128, 142, 113, 125, 111, 132, 133, 109, 107]
alpha = .95

n = len(observations) # find "n"
x_bar = np.mean(observations) # find "x_bar"- the sample mean
sd = np.std(observations) # find the sample standard deviation

t_mult = stats.t.interval(alpha, df = n-1)[1] # Find multiplier

lower = x_bar - t_mult * (sd / np.sqrt(n))
upper = x_bar + t_mult * (sd / np.sqrt(n))
```

```
### BEGIN HIDDEN TESTS
observations = [104, 148, 109, 104, 108, 120, 134, 129, 140, 128, 142, 113, 125, 111, 132, 133, 109, 107]

n = len(observations) # find "n"
x_bar = np.mean(observations) # find "x_bar"- the sample mean
sd = np.std(observations) # find the sample standard deviation
alpha = .95

t_mult = stats.t.interval(alpha, df = n-1)[1]

lower_T = x_bar - t_mult * (sd / np.sqrt(n))
upper_T = x_bar + t_mult * (sd / np.sqrt(n))

assert round(lower_T,3) == round(lower,3)
assert round(upper_T,3) == round(upper,3)

### END HIDDEN TESTS
```

## Question 6

```
### GRADED

### Calculate the 95% confidence interval (with a t-distribution)
### of the data stored in the "observations" variable below

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places

### YOUR ANSWER BELOW

observations = [124, 147, 136, 136, 100, 133, 137, 117, 121, 127, 130, 132, 143,
               146, 130, 149, 119, 146, 107, 148, 125, 105, 116, 130, 117, 117,
               108, 105, 139, 130]

alpha = .95

n = len(observations) # find "n"
x_bar = np.mean(observations) # find "x_bar"- the sample mean
sd = np.std(observations) # find the sample standard deviation

t_mult = stats.t.interval(alpha, df = n-1)[1] # Find multiplier

lower = x_bar - t_mult * (sd / np.sqrt(n))
upper = x_bar + t_mult * (sd / np.sqrt(n))
```

## Question 7

**NOTE:** The below asks for you to calculate the confidence interval using the *z (normal) distribution* instead of the t-distribution

```
### GRADED

### Calculate the 95% confidence interval **(with a Z- (NORMAL) DISTRIBUTION)**
### of the data stored in the "observations" variable below

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places

### YOUR ANSWER BELOW

observations = [124, 147, 136, 136, 100, 133, 137, 117, 121, 127, 130, 132, 143,
               146, 130, 149, 119, 146, 107, 148, 125, 105, 116, 130, 117, 117,
               108, 105, 139, 130]

alpha = .95

n = len(observations) # find "n"
x_bar = np.mean(observations) # find "x_bar"- the sample mean
sd = np.std(observations) # find the sample standard deviation

z_mult = stats.norm.interval(alpha)[1] # Find multiplier

lower = x_bar - z_mult * (sd / np.sqrt(n))
upper = x_bar + z_mult * (sd / np.sqrt(n))
```

## Confidence Intervals: Difference of Means

Let's review the Central Park calculations from lectures 7-5 and 7-6:

Mean temperature between 1869 and 1968 ( $\bar{Y}$ ) is 35.0 with a standard deviation ( $s_y$ ) of 3.8. 100 observations ( $n_y$ )

Mean temperature between 1969 and 2015 ( $\bar{X}$ ) is 38.1 with a standard deviation ( $s_x$ ) of 4.4. 47 observations ( $n_x$ )

**Difference of means:**  $\bar{X} - \bar{Y} = 38.1 - 35.0 = 3.1$

**Standard Error:**  $s.e. = \sqrt{\frac{s_x^2}{n_x} + \frac{s_y^2}{n_y}}$

```
se = np.sqrt(((4.4**2)/47)+((3.8**2)/100))
se
```

To create a confidence interval:

$\bar{X} - \bar{Y} \pm Z_{\alpha/2} \cdot s.e.$

Calculate  $Z_{\alpha/2}$  for the 95% confidence interval:

```
stats.norm.interval(.95)
```

Of note, when the observations are  $n \geq 30$ , the normal distribution may be used for calculating confidence intervals instead of the t-distribution.

```
# Calculate CI
(3.1 + stats.norm.interval(.95)[0]*se, 3.1 + stats.norm.interval(.95)[1]*se)
```

To Calculate the p value:

$1 - P\left(-\frac{\bar{X} - \bar{Y}}{s.e.} < Z < \frac{\bar{X} - \bar{Y}}{s.e.}\right)$   
 $= 2 \cdot P\left(Z < -\frac{3.1}{1.745}\right)$   
 $\approx 2 \cdot P(Z < -4.156)$   
 $\approx .000324$

Calculations below :

```
# Find test statistic using standard error calculated above
test_stat = (38.1-35.0)/ se
print(test_stat)

# Find p-value
print("p-value: ", round((2*(stats.norm.cdf(-test_stat))),7))
```

## Final Example:

One additional example of finding the confidence interval of the difference of two means, and finding the p-value:

```
obs1 = [32.42, 34.61, 35.09, 35.67, 32.04, 34.31, 33.03, 35.55, 34.7 ,
        34.91, 36.02, 32.68, 35.65, 34.14, 32.65, 34.55, 32.78, 37.7 ,
        33.91, 33.53, 31.32, 33.25, 35.07, 36.66, 36.55, 33.52, 33.32,
        32.55, 33.69, 36.05, 30.66, 35.02, 34.05, 34.67, 37.61, 33.71,
        35.72, 34.54, 35.05, 33.69, 30.33, 32.01, 33.16, 36.3 , 32.66,
        31.73, 33.35, 33.16, 33.76, 33.92, 32.05, 34.18, 34.45, 31.49,
        31.9 , 34.33, 33.2 , 31.37, 34.56, 32.61]

obs2 = [36.2 , 41.98, 38.58, 33.59, 36.55, 33.5 , 30.78, 40.87, 42.25,
        39.08, 28.09, 36.74, 44.41, 29.22, 38.55, 24.41, 28.93, 31.97,
        36.6 , 36. , 37.96, 33.92, 43.8 , 36.96, 41.44, 40.54, 35.88,
        30.82, 38.7 , 29.1 ]

# Calculate sample means
mean_x = np.mean(obs1)
mean_y = np.mean(obs2)

# Calculate sample standard deviations
sd_x = np.std(obs1)
sd_y = np.std(obs2)

# Count number of observations in each sample
n_x = len(obs1)
n_y = len(obs2)

# Set alpha
alpha = .95

print("Means: ", mean_x, ", ", mean_y)
print("Standard Deviations: ", sd_x, ", ", sd_y)
print("Number of Observations: ", n_x, ", ", n_y)

# Calculate Observed Difference of means
diff = mean_x - mean_y

# Calculate Standard Error
se = np.sqrt( (sd_x**2/n_x) + (sd_y**2/n_y))

# Find z-multiplier
z_mult = stats.norm.interval(alpha)[1]
```

```

print("\n\nZ-Multiplier: ", z_mult)
print("Difference of Means: ", diff)
print("Standard Error: ", se)

# Calculate confidence interval
lower = diff - z_mult * se
upper = diff + z_mult * se

print("Confidence Interval: ", lower, " , ", upper)

```

With our 95% confidence interval spanning from -3.86 to -0.2 we can say with 95% confidence that the difference in the means of the two sets of observations does not include 0.

Thus we know that the p-value will be less than .05.

Below the exact p-value is calculated

```

# Calculate Test Statistic
test_stat = diff/ se
print(test_stat)

# Find p-value
print("p-value: ", round((2*(stats.norm.cdf(test_stat))),4))

```

Notice, because the difference was negative, the test statistic was also negative, and thus did not need to be made negative when passed to "stats.norm.cdf()"

See below what would happen if it were made negative -- a nonsensical p-value

```

print("WRONG p-value: ", round((2*(stats.norm.cdf(-test_stat))),4))

```

## Question 8

```

### GRADED

### Calculate the 95% confidence interval (with a Z- (NORMAL) DISTRIBUTION)
### of the difference of the means of the collections stored in obs1 and obs2

### NOTE: Specifically find the CI for the mean of obs1 - mean of obs2

### Assign the lower bound as a number to "lower" and the upper bound as a number to "upper"

### Answers will be tested to three decimal places
### YOUR ANSWER BELOW

obs1 = [22.9 , 26.08, 25.04, 22.09, 24.28, 31.3 , 25.47, 24.17, 23.42,
        25.64, 23.96, 23.94, 25.35, 20.92, 27.74, 25.93, 26.9 , 27.87,
        22.43, 23.73, 29.25, 25.66, 23.6 , 26.77, 17.38, 26.26, 17.67,
        24.04, 19.42, 27.41, 30.02, 25.22, 26.47, 24.47, 22.85, 20.07,
        29.46, 23.61, 26.54, 25.37]

obs2 = [26.37, 32.62, 22.13, 22.64, 32.33, 25.62, 18.69, 26.86, 17.87,
        18.16, 26.37, 25.77, 22.57, 27.41, 17.2 , 22.61, 26.97, 28.78,
        24.02, 25.41, 27.88, 28.99, 30.06, 30.23, 24.19, 17.06, 24.38,
        24.13, 25.87, 31.58, 21.19, 32.07, 30.07, 24.23, 27.37]

alpha = .95

# Calculate sample means
mean_x = np.mean(obs1)
mean_y = np.mean(obs2)

# Calculate sample standard deviations
sd_x = np.std(obs1)
sd_y = np.std(obs2)

# Count number of observations in each sample
n_x = len(obs1)
n_y = len(obs2)

# Calculate Observed Difference of means
diff = mean_x - mean_y

# Calculate Standard Error
se = np.sqrt( (sd_x**2/n_x) + (sd_y**2/n_y))

# Find z-multiplier
z_mult = stats.norm.interval(alpha)[1]

# Calculate confidence interval
lower = diff - z_mult * se
upper = diff + z_mult * se

```

## Question 9

```

### GRADED

### Calculate the p-value for the difference of the means of the two samples.

### Answers will be tested to 3 decimal place

### Assign numeric answer to "p_val"
### YOUR ANSWER BELOW

obs1 = [22.9 , 26.08, 25.04, 22.09, 24.28, 31.3 , 25.47, 24.17, 23.42,
        25.64, 23.96, 23.94, 25.35, 20.92, 27.74, 25.93, 26.9 , 27.87,
        22.43, 23.73, 29.25, 25.66, 23.6 , 26.77, 17.38, 26.26, 17.67,
        24.04, 19.42, 27.41, 30.02, 25.22, 26.47, 24.47, 22.85, 20.07,
        29.46, 23.61, 26.54, 25.37]

```

```
obs2 = [26.37, 32.62, 22.13, 22.64, 32.33, 25.62, 18.69, 26.86, 17.87,  
        18.16, 26.37, 25.77, 22.57, 27.41, 17.2 , 22.61, 26.97, 28.78,  
        24.02, 25.41, 27.88, 28.99, 30.06, 30.23, 24.19, 17.06, 24.38,  
        24.13, 25.87, 31.58, 21.19, 32.07, 30.07, 24.23, 27.37]  
  
# Calculate sample means  
mean_x = np.mean(obs1)  
mean_y = np.mean(obs2)  
  
# Calculate sample standard deviations  
sd_x = np.std(obs1)  
sd_y = np.std(obs2)  
  
# Count number of observations in each sample  
n_x = len(obs1)  
n_y = len(obs2)  
  
# Calculate Observed Difference of means  
diff = mean_x - mean_y  
  
# Calculate Standard Error  
se = np.sqrt( (sd_x**2/n_x) + (sd_y**2/n_y))  
  
test_stat = diff/ se  
  
# Find p-value  
p_val = 2*(stats.norm.cdf(test_stat))
```