

Week 2

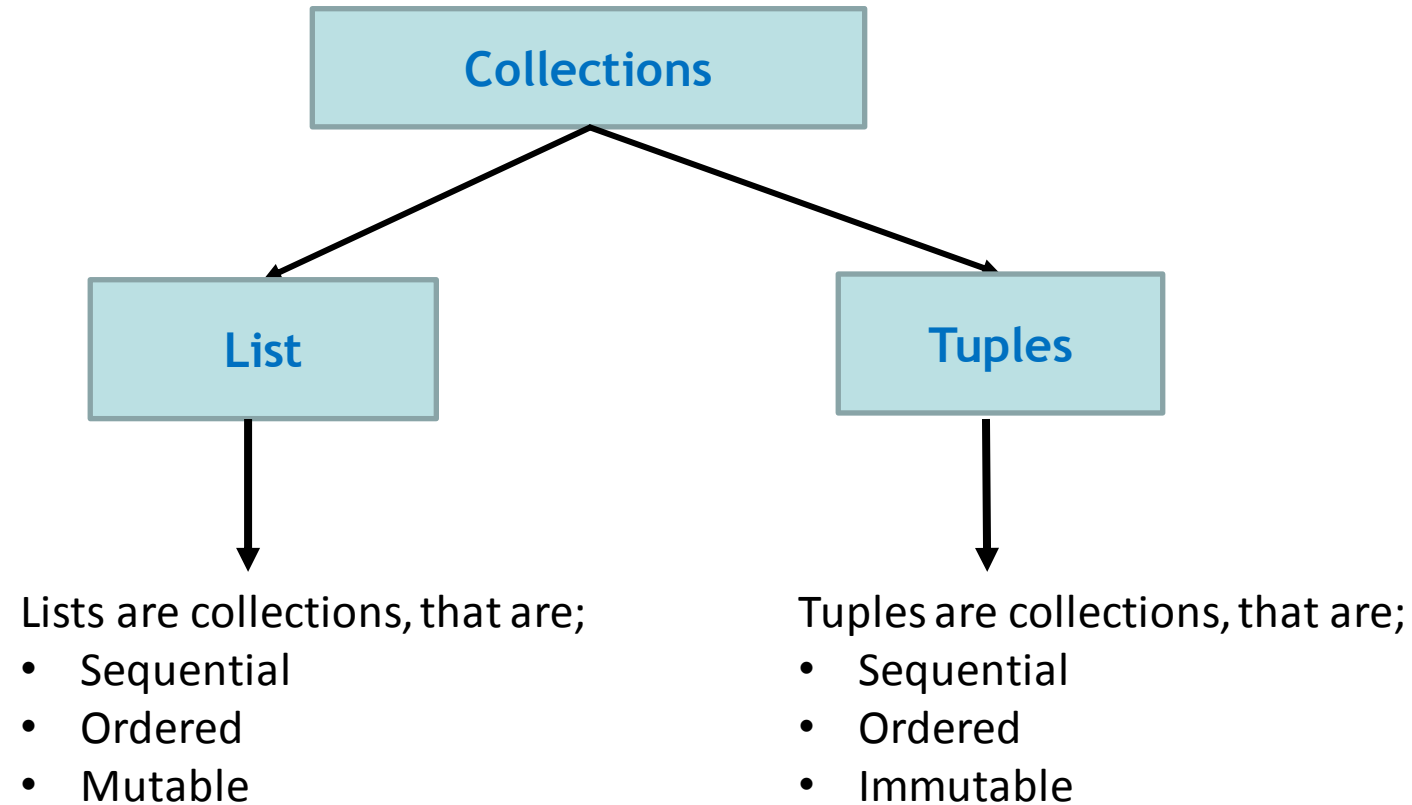
Intermediate Python — Data Structures for your Analysis

Applied Data Science

Columbia University - Columbia Engineering

- ❖ Week 1: Python Basics: How to translate procedures into codes
- ❖ **Week 2: Intermediate Python – Data Structures for Your Analysis**
- ❖ Week 3: Relational Databases – Where Big Data is Typically Stored
- ❖ Week 4: SQL – Ubiquitous Database Format/Language
- ❖ Week 5: Statistical Distributions – The Shape of Data
- ❖ Week 6: Sampling – When You Can't or Won't Have ALL the Data
- ❖ Week 7: Hypothesis Testing – Answering Questions about Your Data
- ❖ Week 8: Data Analysis and Visualization – Using Python's NumPy for Analysis
- ❖ Week 9: Data Analysis and Visualization – Using Python's Pandas for Data Wrangling
- ❖ Week 10: Text Mining – Automatic Understanding of Text
- ❖ Week 11: Machine Learning – Basic Regression and Classification
- ❖ Week 12: Machine learning – Decision Trees and Clustering

In addition to basic data types like strings and floating integers, there are collections.



Key properties

- * collection of related objects
- * ordered or sequential collection
- * mutable. Lists can be modified

Examples

```
list_of_names = ["John", "Jack", "Jill", "Joan"]  
list_of_tickers = ["AAPL", "IONS", "GE", "DB"]  
list_of_natural_numbers = [1, 2, 3, 4, 5, 6, 7]  
long_list = [1, ['a', ['b', 'c']], 43, "Too many cooks spoil the broth"]
```

objects in a list don't
have to be of the same
type



```
long_list = [1,['a',['b','c']],43,"Too many cooks spoil the broth"]
long_list.append('Many hands make light work') #adds an item to the back of the list
long_list[3] #Gets the 4th item in the list
long_list[1][1][0] #Accessing nested items
long_list.extend(['e','f']) #appends contents of a list
long_list.remove(1) #Removes the item with the VALUE 1
long_list.pop() #Removes and returns the last item
long_list.pop(1) #Removes and returns the ith item
len(long_list) #Returns the length of the list
```

Lists: Sequential, Ordered Collection

Creating lists

```
In [ ]: x = [4,2,6,3] #Create a list with values
        y = list() # Create an empty list
        y = [] #Create an empty list
        print(x)
        print(y)
```

Adding items to a list

```
In [ ]: x=list()
        print(x)
        x.append('One') #Adds 'One' to the back of the empty list
        print(x)
        x.append('Two') #Adds 'Two' to the back of the list ['One']
        print(x)
```

```
In [ ]: x.insert(0,'Half') #Inserts 'Half' at location 0. Items will shift to make room
        print(x)
```

```
In [ ]: x=list()
        x.extend([1,2,3]) #Unpacks the list and adds each item to the back of the list
        print(x)
```

Indexing and slicing

```
In [ ]: x=[1,7,2,5,3,5,67,32]
        print(len(x))
        print(x[3])
        print(x[2:5])
        print(x[-1])
        print(x[-1:-1])
```

Removing items from a list

```
In [ ]: x=[1,7,2,5,3,5,67,32]
        x.pop() #Removes the last element from a list
        print(x)
        x.pop(3) #Removes element at item 3 from a list
        print(x)
        x.remove(7) #Removes the first 7 from the list
        print(x)
```

Anything you want to remove must be in the list or the location must be inside the list

```
In [ ]: x.remove(20)
```

Mutability of lists

```
In [ ]: y=['a','b']
        x = [1,7,3]
        print(x)
        print(y)
        y[1] = 4
        print(y)
```

```
In [ ]: print(x)
```

```
In [ ]: x="Hello"
        print(x, id(x))
        x=" You!"
        print(x, id(x)) #x is not the same object it was
        y=["Hello"]
        print(y, id(y))
```


Contents of a list can be changed

Examples

```
x = [1,2,3,4]
x[0]=8 —> [8,2,3,4]
```

immutable: data objects that cannot be changed
e.g. the number 5 is immutable (we can't make it into an 8!)

mutable: data objects that can be changed
e.g., a list of objects owned by Jack and Jill
['pail','water']
(it can be changed to ['pail'])

int, str, bool, float are immutable
list objects are mutable

every python object is either mutable or immutable

Mutable vs Immutable: What's the difference?

```
def eggs(item,total=0):
    total+=item
    return total
print(eggs(1))
print(eggs(2))
```

```
def spam(elem,some_list=[]):
    some_list.append(elem)
    return some_list
print(spam(1))
print(spam(2))
```

```
price = ("20150904", 545.23)
price[0] -> "20140904"
price[1] -> 545.23
price[1]=26.3 -> TypeError
price[2] -> IndexError
```

Tuples are just like lists except they are not mutable
(cannot be changed)

All list operations, except for the ones that change
the value of a list, are also valid tuple operations

Iterating Using Location Indices and Accessing Items Sequentially



`range`: a **sequence** of integers from 0 to length of prices

`len`: the number of items in prices

```
for index in range(len(prices)):
    print(prices[index][0], prices[index][1])
```

`index`: a variable name that holds each value of the sequence in turn. One iteration - one value!

`stock_price`: a variable that will map to each element in the list sequentially

```
prices = [('AAPL', 96.43), ('IONS', 39.28), ('GS', 159.53)]
for stock_price in prices:
    print(stock_price[0], stock_price[1])
```

```
prices = [('AAPL',96.43),('IONS',39.28),('GS',159.53)]
ticker = input("Please enter a ticker: ")
for item in prices:
    if item[0] == ticker:
        print(ticker,item[1])
        break
    else:
        print("Sorry",ticker,"was not found in my database")
print("Statement after for")
```

the for block

else: the program will do this only if the for does not encounter a 'break'

break: the loop will end and control will pass outside the for loop

Range iteration

```
In [ ]: #The for loop creates a new variable (e.g., index below)
#range(len(x)) generates values from 0 to len(x)
x=[1,7,2,5,3,5,67,32]
for index in range(len(x)):
    print(x[index])
```

```
In [ ]: list(range(len(x)))
```

List element iteration

```
In [25]: x=[1,7,2,5,3,5,67,32]
for element in x: #The for draws elements - sequentially - from the list x and uses the variable "element" to store val
    print(element)

1
7
2
5
3
5
67
32
```

Practice problem

Write a function `search_list` that searches a list of tuple pairs and returns the value associated with the first element of the pair

```
In [ ]: def search_list(list_of_tuples,value):
#Write the function here
```

```
In [ ]: prices = [('AAPL',96.43),('IONS',39.28),('GS',159.53)]
ticker = 'IONS'
print(search_list(prices,ticker))
```

Dictionaries

```
In [ ]: mktcaps = {'AAPL':538.7,'GOOG':68.7,'IONS':4.6}
```

```
In [ ]: mktcaps['AAPL'] #Returns the value associated with the key "AAPL"
```

```
In [ ]: mktcaps['GS'] #Error because GS is not in mktcaps
```

Dictionaries are collections that are;

- Unordered
- Pair of elements with a key and a value
- Access values through keys
- Keys are immutable

Sets are collections that are;

- Unordered
- Collection of unique elements
- Does not contain key-value pairs
- Values are immutable

Dictionaries: Key-value Pairs

```
mktcaps = {'AAPL':538.7,'GOOG':68.7,'IONS':4.6}
mktcaps['AAPL'] #key-based retrieval
print(mktcaps['AAPL'])
mktcaps['GE'] #error (no "GE")
'GE' in mktcaps
mktcaps.keys() #returns a list of keys
sorted(mktcaps.keys()) #returns a sorted list of keys
```

Dictionaries

```
In [50]: mktcaps = {'AAPL':538.7,'GOOG':68.7,'IONS':4.6}
         type(mktcaps)
         mktcaps

Out[50]: {'AAPL': 538.7, 'GOOG': 68.7, 'IONS': 4.6}

In [47]: mktcaps['AAPL'] #Returns the value associated with the key 'AAPL'

Out[47]: 538.7

In [48]: mktcaps['GS'] #Error because GS is not in mktcaps
-----
KeyError                                Traceback (most recent call last)
<ipython-input-48-40312e94d4ed> in <module>()
----> 1 mktcaps['GS'] #Error because GS is not in mktcaps

KeyError: 'GS'

In [49]: mktcaps.get('GS') #Returns None because GS is not in mktcaps

In [51]: mktcaps['GS'] = 88.65 #Adds GS to the dictionary
         print(mktcaps)

{'AAPL': 538.7, 'GOOG': 68.7, 'IONS': 4.6, 'GS': 88.65}

In [52]: del(mktcaps['GOOG']) #Removes GOOG from mktcaps
         print(mktcaps)

{'AAPL': 538.7, 'IONS': 4.6, 'GS': 88.65}

In [53]: mktcaps.keys() #Returns all the keys

Out[53]: dict_keys(['AAPL', 'IONS', 'GS'])

In [ ]: mktcaps.values() #Returns all the values

In [ ]: sorted(mktcaps.keys())
```


Sets: Unordered Collections of Unique Objects

```
tickers={"AAPL","GE","NFLX","IONS"}  
regions={"North East","South","West coast","Mid-  
West"}  
"AAPL" in tickers #membership test  
"IBM" not in tickers #non-membership test  
pharma_tickers={"IONS","IMCL"}  
tickers.isdisjoint(pharma_tickers) #empty intersection  
pharma_tickers <= tickers #subset test  
pharma_tickers < tickers #proper-subset test  
tickers > pharma_tickers #superset  
tickers & pharma_tickers #intersection  
tickers | pharma_tickers #union  
tickers - pharma_tickers #set difference
```


In Python, the 'datetime' library is an extremely useful library for data analysis because time is a critical data element.

datetime library

In []:

- Time is linear
- progresses as a straightline trajectory from the big bang
- to now and into the future

Reasoning about time is important in data analysis

- Analyzing financial timeseries data
- Looking at commuter transit passenger flows by time of day
- Understanding web traffic by time of day
- Examining seasonality in department store purchases

The datetime library

- understands the relationship between different points of time
- understands how to do operations on time

Example:

- Which is greater? "10/24/2017" or "11/24/2016"

```
In [ ]: d1 = "10/24/2017"
        d2 = "11/24/2016"
        max(d1,d2)
```

- How much time has passed?

```
In [ ]: d1 - d2
```

Example:

- Which is greater? "10/24/2017" or "11/24/2016"

```
In [4]: d1 = "10/24/2017"
        d2 = "11/24/2016"
        max(d1,d2)
```

```
Out[4]: '11/24/2016'
```

- How much time has passed?

```
In [2]: d1 - d2
```

```
-----
TypeError                                 Traceback (most recent call last)
<ipython-input-2-8e72eafe0703> in <module>()
----> 1 d1 - d2

TypeError: unsupported operand type(s) for -: 'str' and 'str'
```

Obviously that's not going to work.

We can't do date operations on strings

Let's see what happens with datetime

```
In [5]: import datetime
        d1 = datetime.date(2016,11,24)
        d2 = datetime.date(2017,10,24)
        max(d1,d2)
```

```
Out[5]: datetime.date(2017, 10, 24)
```

```
In [ ]: print(d2 - d1)
```

- datetime objects understand time

The Datetime Library: Types 1



The datetime library contains several useful types

- date: stores the date (month,day,year)
- time: stores the time (hours,minutes,seconds)
- datetime: stores the date as well as the time (month,day,year,hours,minutes,seconds)
- timedelta: duration between two datetime or date objects

datetime.date

```
In [7]: import datetime
century_start = datetime.date(2000,1,1)
today = datetime.date.today()
print(century_start,today)
print("We are",today-century_start,"days into this century")
```

```
2000-01-01 2017-05-09
We are 6338 days, 0:00:00 days into this century
```

For a cleaner output

```
In [8]: print("We are",(today-century_start).days,"days into this century")
```

```
We are 6338 days into this century
```

datetime.datetime

```
In [ ]: century_start = datetime.datetime(2000,1,1,0,0,0)
time_now = datetime.datetime.now()
print(century_start,time_now)
print("We are",time_now - century_start,"days, hour, minutes and seconds into this century")
```

datetime objects can check validity

- A ValueError exception is raised if the object is invalid

datetime.timedelta

Used to store the duration between two points in time

```
In [15]: century_start = datetime.datetime(2000,1,1,0,0,0)
time_now = datetime.datetime.now()
time_since_century_start = time_now - century_start
print("days since century start",time_since_century_start.days)
print("seconds since century start",time_since_century_start.total_seconds())
print("minutes since century start",time_since_century_start.total_seconds()/60)
print("hours since century start",time_since_century_start.total_seconds()/60/60)
```

```
days since century start 6338
seconds since century start 547640865.589801
minutes since century start 9127347.759830017
hours since century start 152122.4626638336
```

datetime.time

```
In [16]: date_and_time_now = datetime.datetime.now()
time_now = date_and_time_now.time()
print(time_now)
```

```
10:28:06.130552
```

You can do arithmetic operations on datetime objects

- You can use timedelta objects to calculate new dates or times from a given date

```
In [ ]: today=datetime.date.today()
five_days_later=today+datetime.timedelta(days=5)
print(five_days_later)
```

```
In [ ]: now=datetime.datetime.today()
five_minutes_and_five_seconds_later = now + datetime.timedelta(minutes=5,seconds=5)
print(five_minutes_and_five_seconds_later)
```

```
In [ ]: now=datetime.datetime.today()
five_minutes_and_five_seconds_earlier = now-datetime.timedelta(minutes=5,seconds=5)
print(five_minutes_and_five_seconds_earlier)
```

- But you can't use timedelta on time objects. If you do, you'll get a TypeError exception

```
In [ ]: time_now=datetime.datetime.now().time() #Returns the time component (drops the day)
```

The Datetime Library: Types 2

NAME

strptime - date and time conversion

SYNOPSIS

```
1 #include <time.h>
char *strptime(const char *restrict buf, const char *restrict format,
               struct tm *restrict tm);
```

DESCRIPTION

The `strptime()` function shall convert the character string pointed to by `buf` to values which are stored in the `tm` structure pointed to by `tm`, using the format specified by `format`.

The `format` is composed of zero or more directives. Each directive is composed of one of the following: one or more white-space characters (as specified by `isspace()`); an ordinary character (neither `'%'` nor a white-space character); or a conversion specification. Each conversion specification is composed of a `'%'` character followed by a conversion character which specifies the replacement required. The application shall ensure that there is white-space or other non-alphanumeric characters between any two conversion specifications. The following conversion specifications are supported:

<code>%a</code>	The day of the week, using the locale's weekday names; either the abbreviated or full name may be specified.
<code>%A</code>	Equivalent to <code>%a</code> .
<code>%b</code>	The month, using the locale's month names; either the abbreviated or full name may be specified.
<code>%B</code>	Equivalent to <code>%b</code> .
<code>%c</code>	Replaced by the locale's appropriate date and time representation.
<code>%C</code>	The century number [00,99]; leading zeros are permitted but not required.
<code>%d</code>	The day of the month [01,31]; leading zeros are permitted but not required.
<code>%D</code>	The date as <code>%m / %d / %y</code> .
<code>%e</code>	Equivalent to <code>%d</code> .
<code>%h</code>	Equivalent to <code>%b</code> .
<code>%H</code>	The hour (24-hour clock) [00,23]; leading zeros are permitted but not required.
<code>%I</code>	The hour (12-hour clock) [01,12]; leading zeros are permitted but not required.
<code>%j</code>	The day number of the year [001,366]; leading zeros are permitted but not required.
<code>%m</code>	The month number [01,12]; leading zeros are permitted but not required.
<code>%M</code>	The minute [00,59]; leading zeros are permitted but not required.
<code>%s</code>	Any white space.
<code>%p</code>	The locale's equivalent of a.m or p.m.
<code>%r</code>	12-hour clock time using the AM/PM notation if <code>t_fmt_ampm</code> is not an empty string in the <code>LC_TIME</code> portion of the current locale; in the POSIX locale, this shall be equivalent to <code>%I : %M : %S %p</code> .

datetime and strings

More often than not, the program will need to get the date or time from a string:
From a website (bus/train timings)
From a file (date or datetime associated with a stock price)
From the user (from the input statement)

Python needs to parse the string so that it correctly creates a date or time object

datetime.strptime

- `datetime.strptime()`: grabs time from a string and creates a date or datetime or time object
- The programmer needs to tell the function what format the string is using
- See <http://pubs.opengroup.org/onlinepubs/009895399/functions/strptime.html> for how to specify the format

```
In [ ]: date='01-Apr-03'
date_object=datetime.datetime.strptime(date,'%d-%b-%y')
print(date_object)
```

```
In [ ]: #Unfortunately, there is no similar thing for time deltas
#So we have to be creative!
bus_travel_time='2:15:30'
hours,minutes,seconds=bus_travel_time.split(':')
x=datetime.timedelta(hours=int(hours),minutes=int(minutes),seconds=int(seconds))
print(x)
```

```
In [ ]: #Or write a function that will do this for a particular format
def get_timedelta(time_string):
    hours,minutes,seconds = time_string.split(':')
    import datetime
    return datetime.timedelta(hours=int(hours),minutes=int(minutes),seconds=int(seconds))
```


Bucketing Time Part 1

Bucketing time

The file "sample_data.csv" contains start times and processing times for all complaints registered with New York City's 311 complaint hotline on 01/01/2016. Our goal is to compute the average processing time for each hourly bucket.

Let's take a quick look at the data

```
In [1]: #Unfortunately, this won't work on Windows.
        !head sample_data.csv
```

```
2016-01-01 00:00:09,0.0815162037037037
2016-01-01 00:00:40,0.1334837962962963
2016-01-01 00:01:09,20.388726851851853
2016-01-01 00:02:59,0.9811458333333334
2016-01-01 00:03:03,7.048576388888889
2016-01-01 00:03:03,0.1400810185185185
2016-01-01 00:03:29,0.11086805555555555
2016-01-01 00:04:06,0.016967592592592593
2016-01-01 00:04:37,0.15972222222222222
2016-01-01 00:04:56,2.996585648148148
```

Step 1: Read the data

```
In [ ]: data_tuples = list()
        with open('sample_data.csv','r') as f:
            for line in f:
                data_tuples.append(line.strip().split(','))
```

Let's look at the first 10 lines

```
In [ ]: data_tuples[0:10]
```

- Element 1 of the tuple is a date inside a string
- Element 2 is double inside a string
- Let's convert them

```
In [6]: #Figure out the format string
        # http://pubs.opengroup.org/onlinepubs/009695399/functions/strftime.html
        import datetime
        x='2016-01-01 00:00:09'
        format_str = '%Y-%m-%d %H:%M:%S'
        datetime.datetime.strptime(x,format_str)
```

```
Out[6]: datetime.datetime(2016, 1, 1, 0, 0, 9)
```

```
In [7]: data_tuples = list()
        with open('sample_data.csv','r') as f:
            for line in f:
                data_tuples.append(line.strip().split(','))
        import datetime
        for i in range(0,len(data_tuples)):
            data_tuples[i][0] = datetime.datetime.strptime(data_tuples[i][0],format_str)
            data_tuples[i][1] = float(data_tuples[i][1])
```

```
In [8]: #Let's see if this worked
        data_tuples[0:10]
```

```
Out[8]: [(datetime.datetime(2016, 1, 1, 0, 0, 9), 0.0815162037037037),
          (datetime.datetime(2016, 1, 1, 0, 0, 40), 0.1334837962962963),
          (datetime.datetime(2016, 1, 1, 0, 1, 9), 20.388726851851853),
          (datetime.datetime(2016, 1, 1, 0, 2, 59), 0.9811458333333334),
          (datetime.datetime(2016, 1, 1, 0, 3, 3), 7.048576388888889),
          (datetime.datetime(2016, 1, 1, 0, 3, 3), 0.1400810185185185),
          (datetime.datetime(2016, 1, 1, 0, 3, 29), 0.11086805555555555),
          (datetime.datetime(2016, 1, 1, 0, 4, 6), 0.016967592592592593),
          (datetime.datetime(2016, 1, 1, 0, 4, 37), 0.15972222222222222),
          (datetime.datetime(2016, 1, 1, 0, 4, 56), 2.996585648148148)]
```

We can replace the datetime by hourly buckets

```
In [ ]: #Extract the hour from a datetime object
        x=data_tuples[0][0]
        x.hour
```

Bucketing Time Part 2



File Edit View Insert Cell Kernel Widgets Help Python 3

We can replace the datetime by hourly buckets

```
In [9]: #Extract the hour from a datetime object
x=data_tuples[0][0]
x.hour

Out[9]: 0
```

Use list comprehension to bucket the data

```
In [10]: for x in data_tuples:
          print(x[0].hour,x[1])

0 0.10398148148148148
0 0.2975231481481482
0 0.09293981481481481
0 0.016446759259259258
0 0.06824074074074074
0 0.04800925925925926
0 0.26761574074074074
0 1.4127662037037036
0 0.4363078703703704
0 0.869375
0 0.05337962962962963
0 0.6558333333333334
0 0.3119560185185185
0 3.580636574074074
0 0.1267939814814815
0 5.040613425925926
0 0.022662037037037036
0 0.31908564814814818
0 1.001412037037037
0 0.3857407407407407
```

```
In [11]: data_tuples = [(x[0].hour,x[1]) for x in data_tuples]
```

```
In [ ]: |
```

```
In [ ]: data_tuples = list()
with open('sample_data.csv','r') as f:
    for line in f:
        data_tuples.append(line.strip().split(','))
import datetime
for i in range(0,len(data_tuples)):
    data_tuples[i][0] = datetime.datetime.strptime(data_tuples[i][0],format_str)
    data_tuples[i][1] = float(data_tuples[i][1])
```

Create a function that returns the data

```
In [14]: def get_data(filename):
          data_tuples = list()
          with open(filename,'r') as f:
              for line in f:
                  data_tuples.append(line.strip().split(','))
          import datetime
          format_str = "%Y-%m-%d %H:%M:%S"
          data_tuples = [(datetime.datetime.strptime(x[0],format_str).hour,float(x[1])) for x in data_tuples]
          return data_tuples
```

```
In [15]: get_data('sample_data.csv')
```

```
(0, 0.059479166666666666),
(0, 0.003460648148148148),
(0, 0.22096064814814814),
(0, 0.10398148148148148),
(0, 0.2975231481481482),
(0, 0.09293981481481481),
(0, 0.016446759259259258),
(0, 0.06824074074074074),
(0, 0.04800925925925926),
(0, 0.26761574074074074),
(0, 1.4127662037037036),
(0, 0.4363078703703704),
(0, 0.869375),
(0, 0.05337962962962963),
(0, 0.6558333333333334),
(0, 0.3119560185185185),
(0, 3.580636574074074),
(0, 0.1267939814814815),
(0, 5.040613425925926),
(0, 0.022662037037037036),
```


Bucketing Time Part 3



Let's print them to see what sort of pattern is there in the data

Bear in mind that this is just one day's data!

```
In [19]: for key,value in buckets.items():  
        print("Hour:",key,"\\tAverage:",value[1]/value[0])
```

Hour: 0	Average: 0.6570511564469035
Hour: 1	Average: 2.9613477328431377
Hour: 2	Average: 2.334965452261305
Hour: 3	Average: 3.0839338759007866
Hour: 4	Average: 4.663183017810805
Hour: 5	Average: 2.550054976851854
Hour: 6	Average: 5.344349026388891
Hour: 7	Average: 2.5844597678664565
Hour: 8	Average: 6.0724520669659565
Hour: 9	Average: 8.564869090207626
Hour: 10	Average: 12.671294691132733
Hour: 11	Average: 5.901653566341063
Hour: 12	Average: 13.66402543540564
Hour: 13	Average: 8.593492462013293
Hour: 14	Average: 8.100135135135135
Hour: 15	Average: 12.776634463154863
Hour: 16	Average: 10.943701434277418
Hour: 17	Average: 6.634365784623489
Hour: 18	Average: 7.324956692612944
Hour: 19	Average: 9.098796085858586
Hour: 20	Average: 5.199433822667603
Hour: 21	Average: 4.74319171267541
Hour: 22	Average: 8.449229102956167
Hour: 23	Average: 5.184938602292768

Step 2: Accumulate counts and sums for each bucket

```
In [16]: buckets = dict()  
for item in get_data('sample_data.csv'):  
    if item[0] in buckets:  
        buckets[item[0]][0] += 1  
        buckets[item[0]][1] += item[1]  
    else:  
        buckets[item[0]] = [1,item[1]]
```

```
In [17]: buckets
```

```
Out[17]: {0: [241, 158.34932870370175],  
1: [340, 1006.8582291666668],  
2: [199, 464.65812499999997],  
3: [221, 681.5493865740739],  
4: [157, 732.1197337962964],  
5: [112, 285.60615740740764],  
6: [80, 427.54798611111124],  
7: [71, 183.4966435185184],  
8: [99, 601.1727546296297],  
9: [132, 1130.5627199074067],  
10: [117, 1735.9673726851845],  
11: [182, 1074.1009490740735],  
12: [168, 2295.5562731481473],  
13: [195, 1675.7310300925922],  
14: [185, 1498.5249999999999],  
15: [193, 2465.8904513888889],  
16: [204, 2232.515092592593],  
17: [211, 1399.8511805555556],  
18: [182, 1333.1421180555558],  
19: [165, 1501.3013541666667],  
20: [158, 821.5105439814813],  
21: [161, 763.653865740741],  
22: [218, 1841.9319444444443],  
23: [210, 1088.8371064814814]}
```

Bucketing Time Part 3



Put everything into a function

This way, we can easily test other similar datasets

```
In [20]: def get_hour_bucket_averages(filename):  
    def get_data(filename):  
        data_tuples = list()  
        with open(filename, 'r') as f:  
            for line in f:  
                data_tuples.append(line.strip().split(','))  
        import datetime  
        format_str = "%Y-%m-%d %H:%M:%S"  
        data_tuples = [(datetime.datetime.strptime(x[0], format_str).hour, float(x[1])) for x in data_tuples]  
        return data_tuples  
    buckets = dict()  
    for item in get_data(filename):  
        if item[0] in buckets:  
            buckets[item[0]][0] += 1  
            buckets[item[0]][1] += item[1]  
        else:  
            buckets[item[0]] = [1, item[1]]  
    return [(key, value[1]/value[0]) for key, value in buckets.items()]
```

```
In [21]: get_hour_bucket_averages('sample_data.csv')
```

```
Out[21]: [(0, 0.6570511564469035),  
(1, 2.9613477328431377),  
(2, 2.334965452261305),  
(3, 3.0839338759007866),  
(4, 4.663183017810805),  
(5, 2.550054976851854),  
(6, 5.344349826388891),  
(7, 2.5844597678664565),  
(8, 6.0724520669639565),  
(9, 8.564869090207626),
```

The file `all_data.csv` contains data from January to September 2016

We can test whether our one day result is generally true or not

```
In [2]: get_hour_bucket_averages('all_data.csv')
```

```
Out[2]: [(0, 4.485612099128487),  
(1, 2.8263083049680278),  
(2, 2.859209391496003),  
(3, 2.9813212672915657),  
(4, 3.520777693173893),  
(5, 4.028842839550067),  
(6, 5.3501635019789914),  
(7, 4.305984716000046),  
(8, 5.090230597495249),  
(9, 6.767684356105564),  
(10, 7.252764762298842),  
(11, 7.156706204701707),  
(12, 7.422673351052525),  
(13, 7.402425948682307),  
(14, 7.546603227374128),  
(15, 8.001251635520441),  
(16, 8.191847429766709),  
(17, 7.275740883284791),  
(18, 6.464817194100053),  
(19, 5.6403138675375155),  
(20, 4.989414785443646),  
(21, 4.275270320395889),  
(22, 3.5846441619204086),  
(23, 3.0346464768596855)]
```

