



EMERITUS
INSTITUTE OF MANAGEMENT



WEEK 6

SQL — UBIQUITOUS DATABASE FORMAT/LANGUAGE



MIT Sloan
Executive Education



Columbia Business School
AT THE VERY CENTER OF BUSINESS™
| EXECUTIVE EDUCATION



Tuck
EXECUTIVE EDUCATION
at Dartmouth

SQL: Structured Query Language and its Elements

➡ The language for relational databases

➡ SQL is a *declarative language*

➡ SQL queries act on tables

➡ SQL queries return tables

- A **table** is a logically connected set of data organized as a set of columns and rows

A **record** is a row of information about a single item in a table

- A **field** is an individual data item in a record
- Each field has a precise **format** (e.g., number, string, date, etc.)
- A **key** field is a field whose value uniquely identifies a record in a table

SQL Servers

- ➡ The program that manages the database
- ➡ Provides database services to client programs
- ➡ Usually modeled as a **client server** system
- ➡ Numerous choices:
 - ➡ Microsoft SQL Server
 - ➡ Oracle
 - ➡ DB2
 - ➡ PostgreSQL
 - ➡ MySQL

MySQL

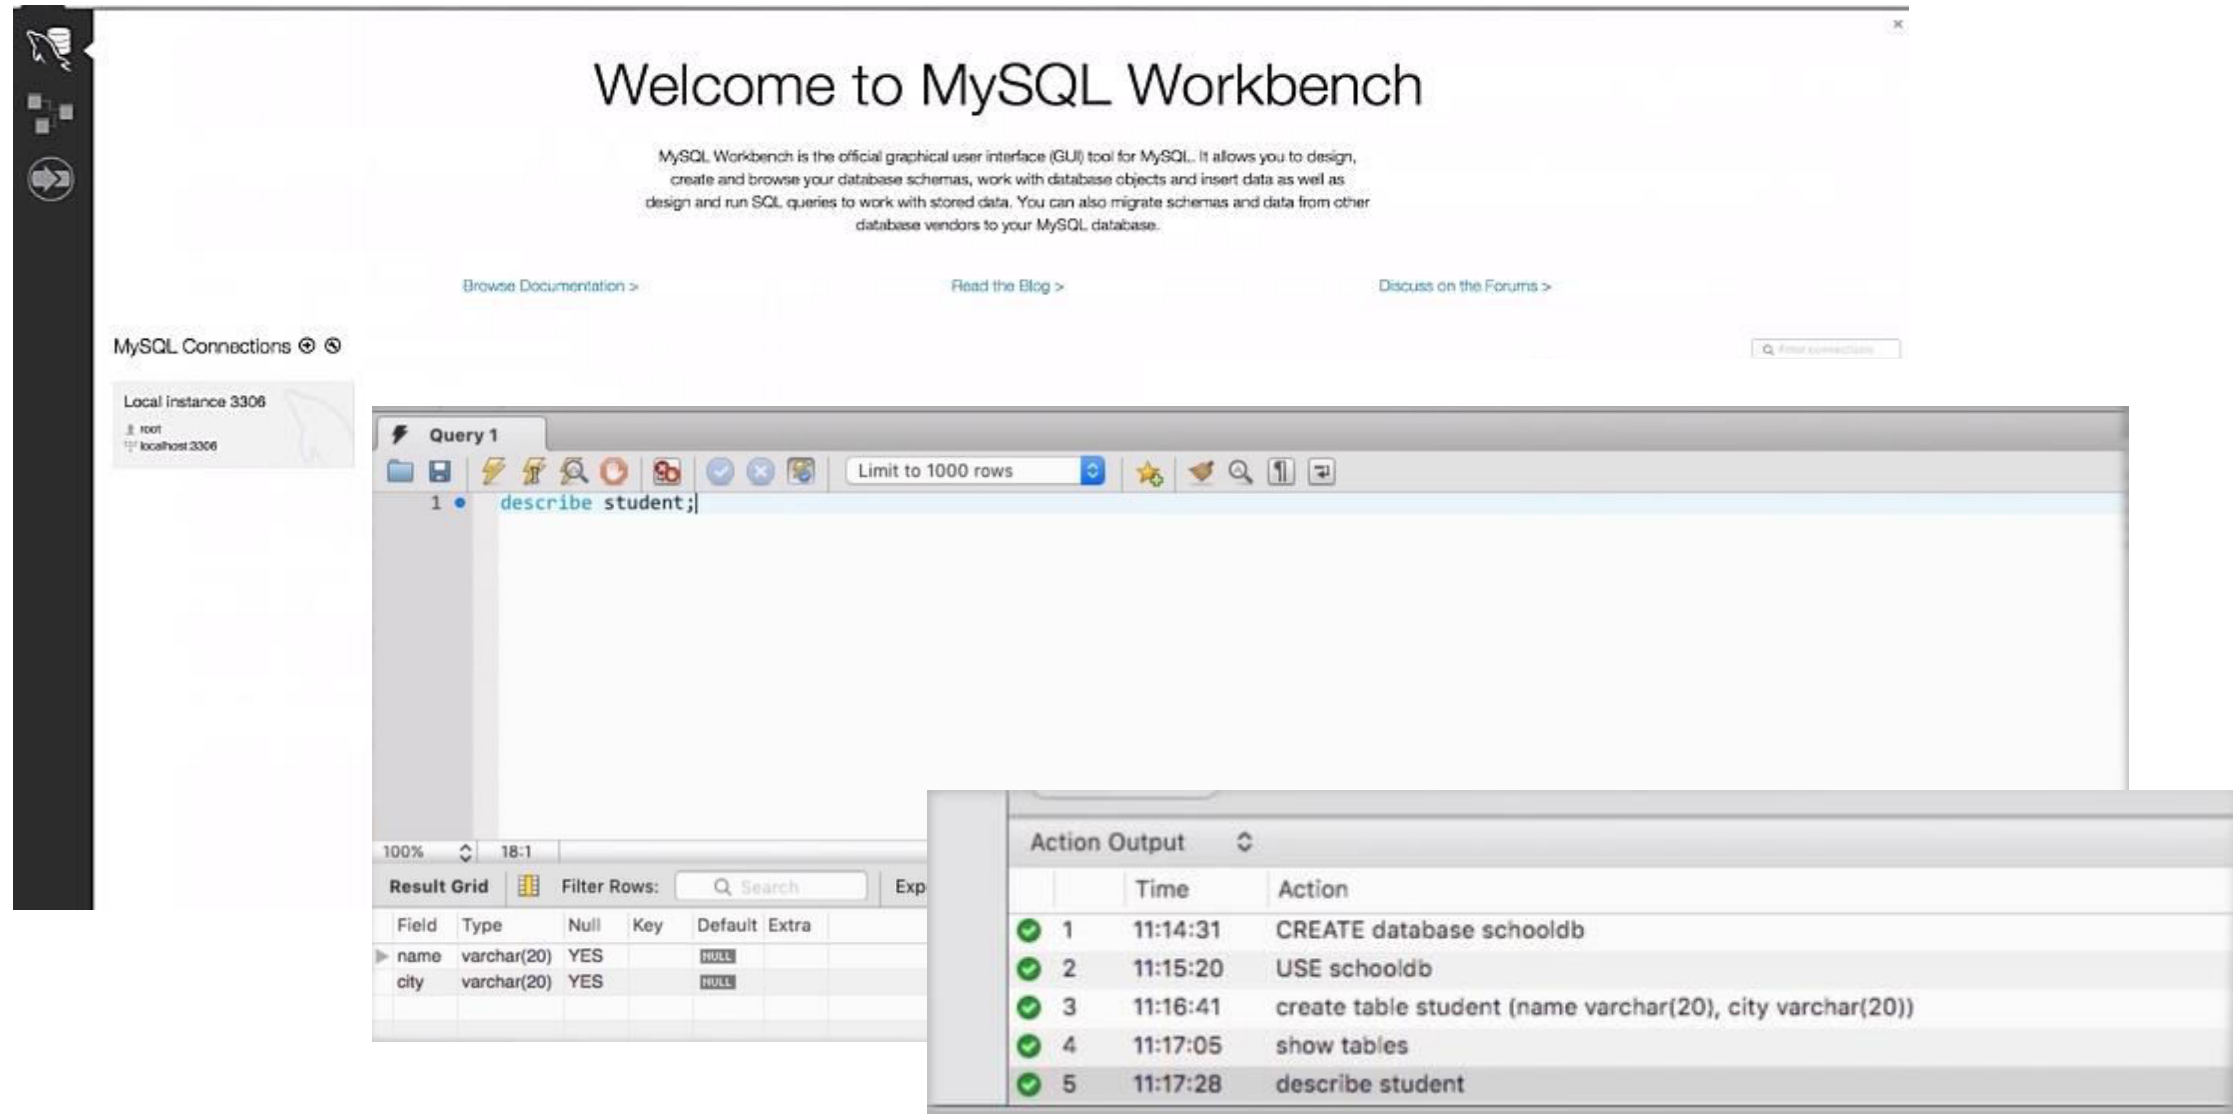
- An open source relational DBMS
- MySQL server is where the database resides
- MySQL clients are the programs that talk to the server
- One server can have many clients

Query Types

- Data definitional
create, alter, drop
- Data manipulation
insert, delete, update
- Data querying
select
select where
select join
- List existing databases:
 - `SHOW databases;`
- Create new database:
 - `CREATE database IF NOT EXISTS SchoolDB;`
- Choose a DB to work on
 - `USE SchoolDB;`
- Delete a database:
 - `DROP DATABASE SchoolDB;`

Introduction to MySQL Workbench

MySQL Workbench is a GUI client for MySQL



Example Database

Student

ssn	f name	l name	phone	city	zip
111-22-3333	John	Childs	646.123.1212	New York	10025
123-12-1234	Mary	Arias		New York	10011
555-11-7777	Roberto	Perez	917.333.5479	San Francisco	94110
222-33-4455	Lila	Pennington	425.123.1212	Seattle	98105



Course

number	name	room
c1	Data analytics	1127
c2	Python	303
c3	corp fin	331
c4	prod. mgmt	1127
c5	Ethics	303
c6	leadership	303
c7	bus analytics	1127

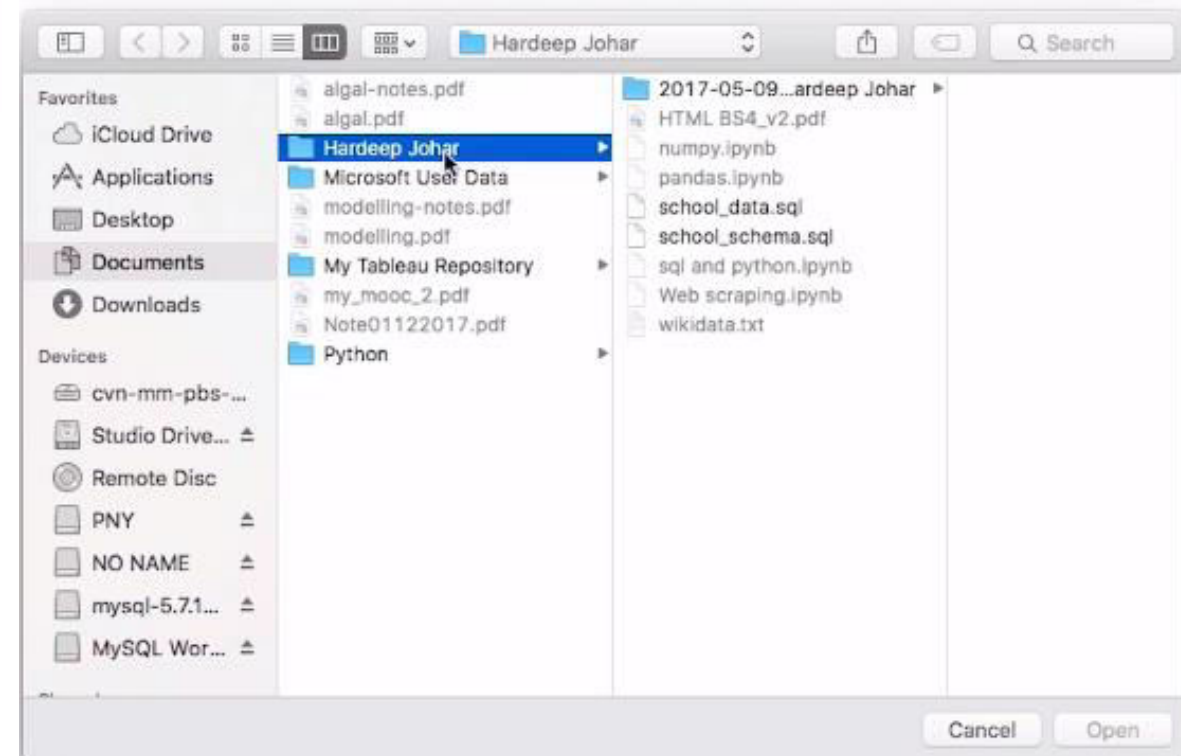
Enrolls-in

ssn	class	score
111-22-3333	c1	93
123-12-1234	c1	87
111-22-3333	c2	95
222-33-4455	c1	44
555-11-7777	c1	36

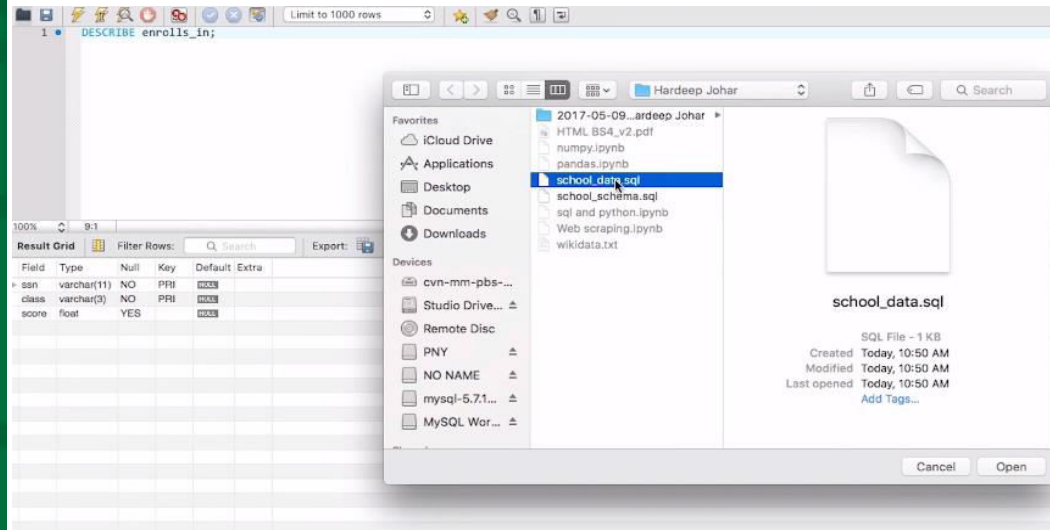
Setup the Database

- Download the files
 - school_schema.sql
 - school_data.sql
- Open mysql workbench (see installation instructions)
- Open school_schema.sql (click the  icon)
- Execute the script (click the  icon)
- Repeat with the file school_data.sql

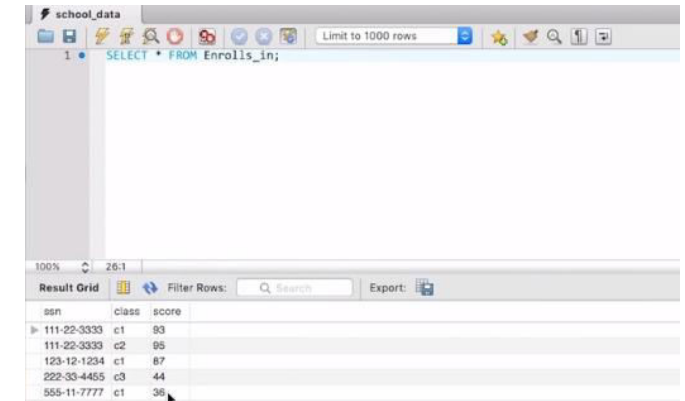
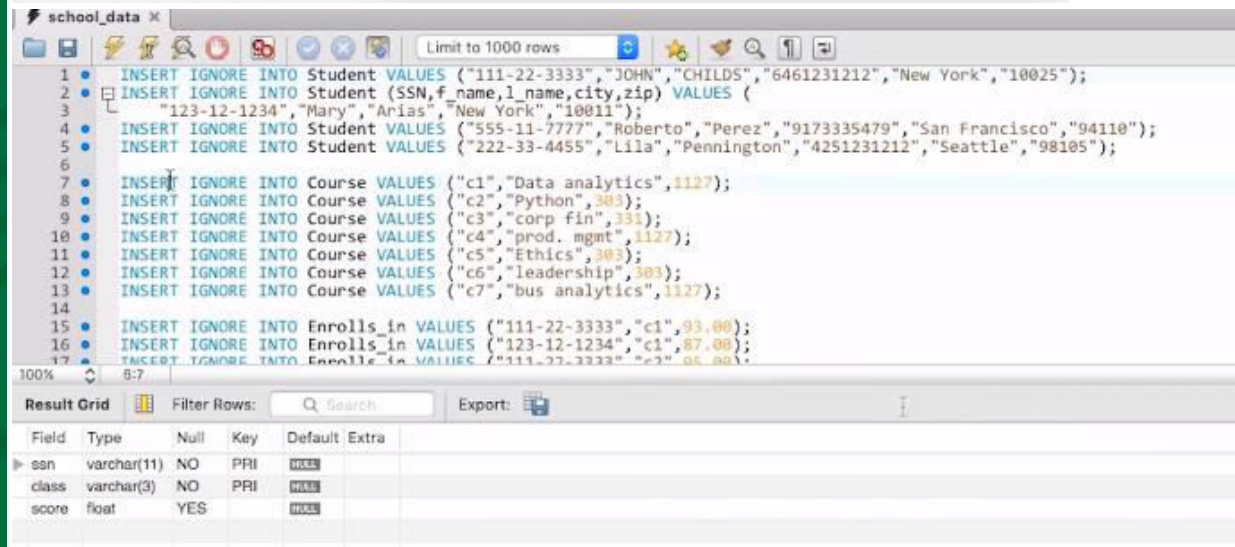
```
1 • DROP DATABASE IF EXISTS schooldb;
2 • CREATE DATABASE IF NOT EXISTS schooldb;
3 • USE schooldb;
4 • CREATE TABLE IF NOT EXISTS Student (
5     ssn VARCHAR(11) NOT NULL PRIMARY KEY,
6     f_name VARCHAR(20) NOT NULL,
7     l_name VARCHAR(20) NOT NULL,
8     phone VARCHAR(10),
9     city VARCHAR(20) NOT NULL,
10    zip VARCHAR(5) NOT NULL
11 );
12
13 • CREATE TABLE IF NOT EXISTS Course (
14     number VARCHAR(3) NOT NULL PRIMARY KEY,
15     name VARCHAR(30) NOT NULL,
16     room INT NOT NULL
17 );
18
19 • CREATE TABLE IF NOT EXISTS Enrolls_in (
20     ssn VARCHAR(11) NOT NULL,
21     class VARCHAR(3) NOT NULL,
22     score FLOAT,
23     CONSTRAINT pk_enroll PRIMARY KEY (ssn,class)
24 );
```



Database Example



```
1 • INSERT IGNORE INTO Student VALUES ("111-22-3333","JOHN","CHILDS","6461231212","New York","10025");
2 • INSERT IGNORE INTO Student (SSN,f_name,l_name,city,zip) VALUES (
3 •   "123-12-1234","Mary","Arias","New York","10011");
4 • INSERT IGNORE INTO Student VALUES ("555-11-7777","Roberto","Perez","9173335479","San Francisco","94110");
5 • INSERT IGNORE INTO Student VALUES ("222-33-4455","Lila","Pennington","4251231212","Seattle","98105");
6
7 • INSERT IGNORE INTO Course VALUES ("c1","Data analytics",1127);
8 • INSERT IGNORE INTO Course VALUES ("c2","Python",303);
9 • INSERT IGNORE INTO Course VALUES ("c3","corp fin",331);
10 • INSERT IGNORE INTO Course VALUES ("c4","prod. mgmt",1127);
11 • INSERT IGNORE INTO Course VALUES ("c5","Ethics",303);
12 • INSERT IGNORE INTO Course VALUES ("c6","leadership",303);
13 • INSERT IGNORE INTO Course VALUES ("c7","bus analytics",1127);
14
15 • INSERT IGNORE INTO Enrolls_in VALUES ("111-22-3333","c1",93.00);
16 • INSERT IGNORE INTO Enrolls_in VALUES ("123-12-1234","c1",87.00);
17 • INSERT IGNORE INTO Enrolls_in VALUES ("111-22-3333","c2",95.00);
18 • INSERT IGNORE INTO Enrolls_in VALUES ("222-33-4455","c3",44.00);
19 • INSERT IGNORE INTO Enrolls_in VALUES ("555-11-7777","c1",36.00);
```




The Select Statement Part 1

- SELECT returns a set of records from one or more tables in a database
- SELECT describes the result set - processing is left to the server

Select

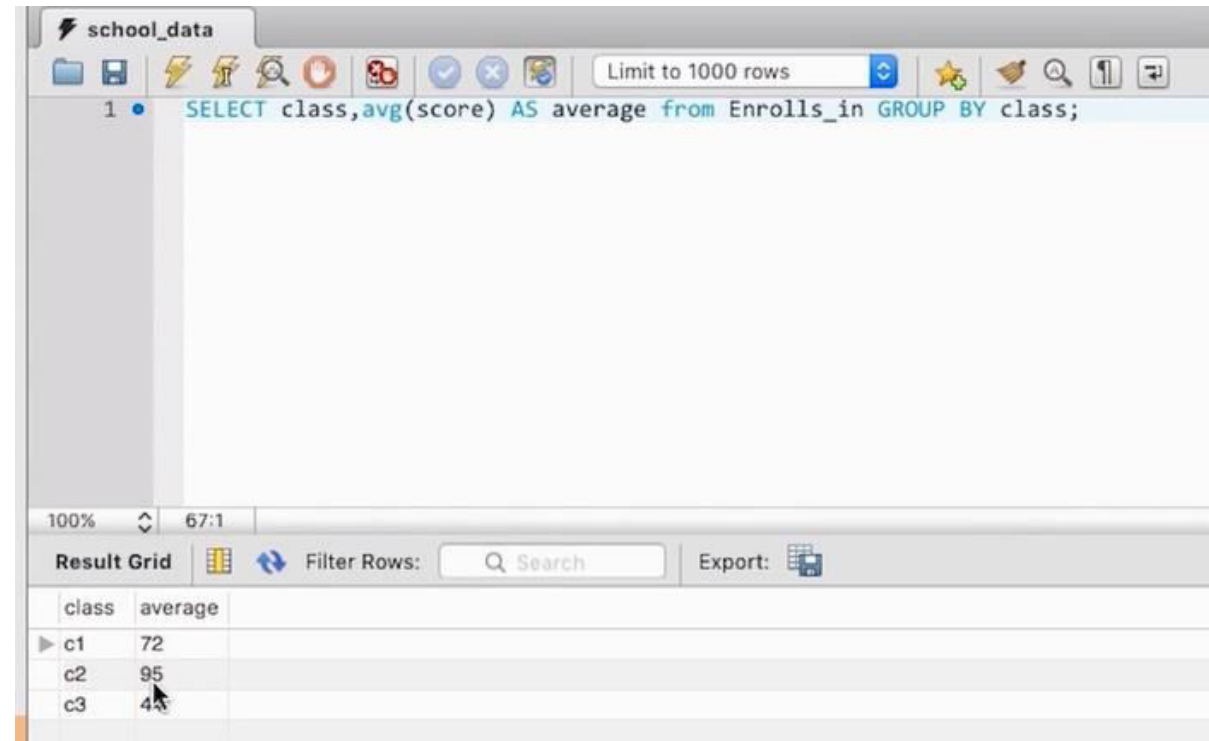
1. Get everything in a table
`SELECT * FROM Student;`
2. Get all students from New York
`SELECT * FROM Student WHERE City = "New York";`
3. Get only the names of students from New York
`SELECT f_name, l_name FROM Student WHERE city = "New York";`
4. Get the ssn of students ordered by f_name
`SELECT ssn FROM Student ORDER BY f_name ASC;`
`SELECT ssn FROM Student ORDER BY f_name DESC;`
5. List the set of unique cities in the Student table
`SELECT DISTINCT city FROM Student;`



The Select Statement Part 2

Select

1. Get the average score for all students
`SELECT avg(score) FROM Enrolls_in;`
2. Get the average score in class c1
`SELECT avg(score) FROM Enrolls_in WHERE class = "c1";`
3. Get the average score for each class
`SELECT class, avg(score) FROM Enrolls_in GROUP BY class;`
4. Get the average score for each class with more than one student
`SELECT class, avg(score) FROM Enrolls_in GROUP BY class
HAVING count(score) > 1;`



The screenshot shows a database query tool interface. At the top, a tab labeled 'school_data' is active. Below the tab, a toolbar contains various icons for file operations, execution, and viewing. A text box shows the SQL query: `1 • SELECT class, avg(score) AS average from Enrolls_in GROUP BY class;`. Below the query, a 'Result Grid' section displays the results of the query. The grid has two columns: 'class' and 'average'. The results are as follows:

class	average
c1	72
c2	95
c3	45

The Select Statement Part 3

Select

1. Get the average score for each class and name the column average

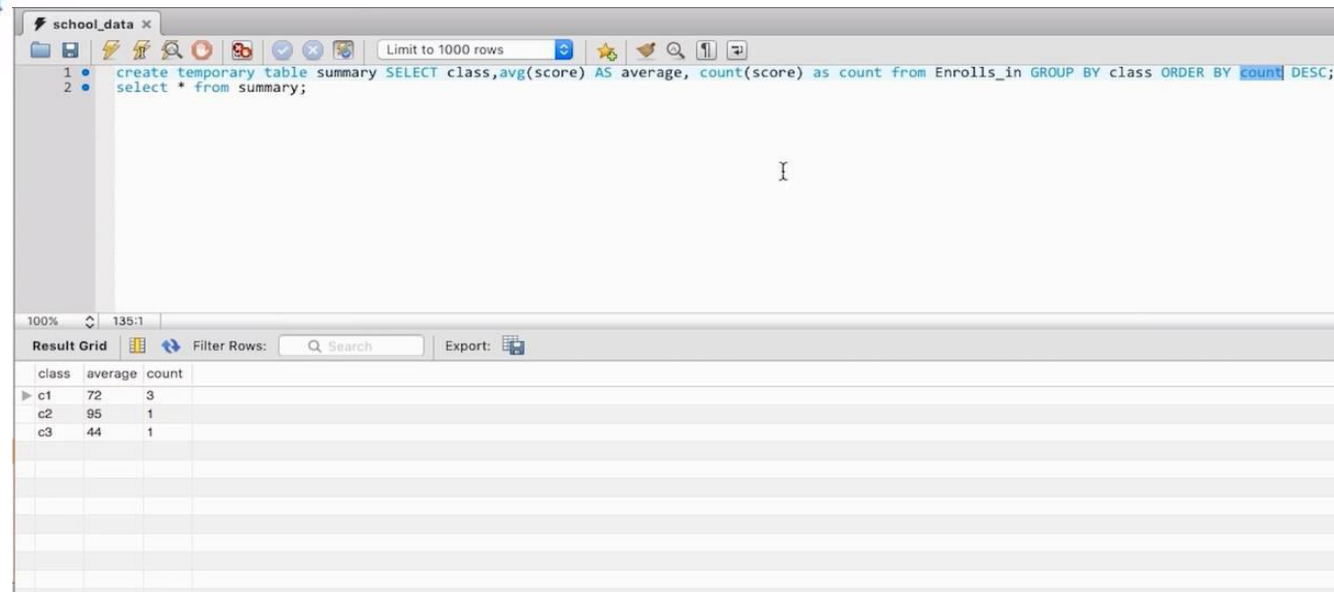
```
SELECT class, avg(score) AS average FROM Enrolls_in  
GROUP BY class;
```

2. Get the average score for each class, name the column average, and store the result set in a temporary table "averages"

```
CREATE TEMPORARY TABLE averages SELECT class,  
avg(score) AS average FROM Enrolls_in GROUP BY class;
```

3. Get the average score and number of students for each class, name the columns average and count, sort them in ascending order of count, and store the result set in a temporary table "summary"

```
CREATE TEMPORARY TABLE IF NOT EXISTS summary  
SELECT class, avg(score) AS average, count(score) AS count  
FROM Enrolls_in GROUP BY class ORDER BY count;
```



The screenshot shows a SQL IDE window titled "school_data". The query window contains the following SQL code:

```
1 • create temporary table summary SELECT class, avg(score) AS average, count(score) as count from Enrolls_in GROUP BY class ORDER BY count DESC;  
2 • select * from summary;
```

The result grid below the query window displays the following data:

class	average	count
c1	72	3
c2	95	1
c3	44	1

Working Across Multiple Tables Part 1

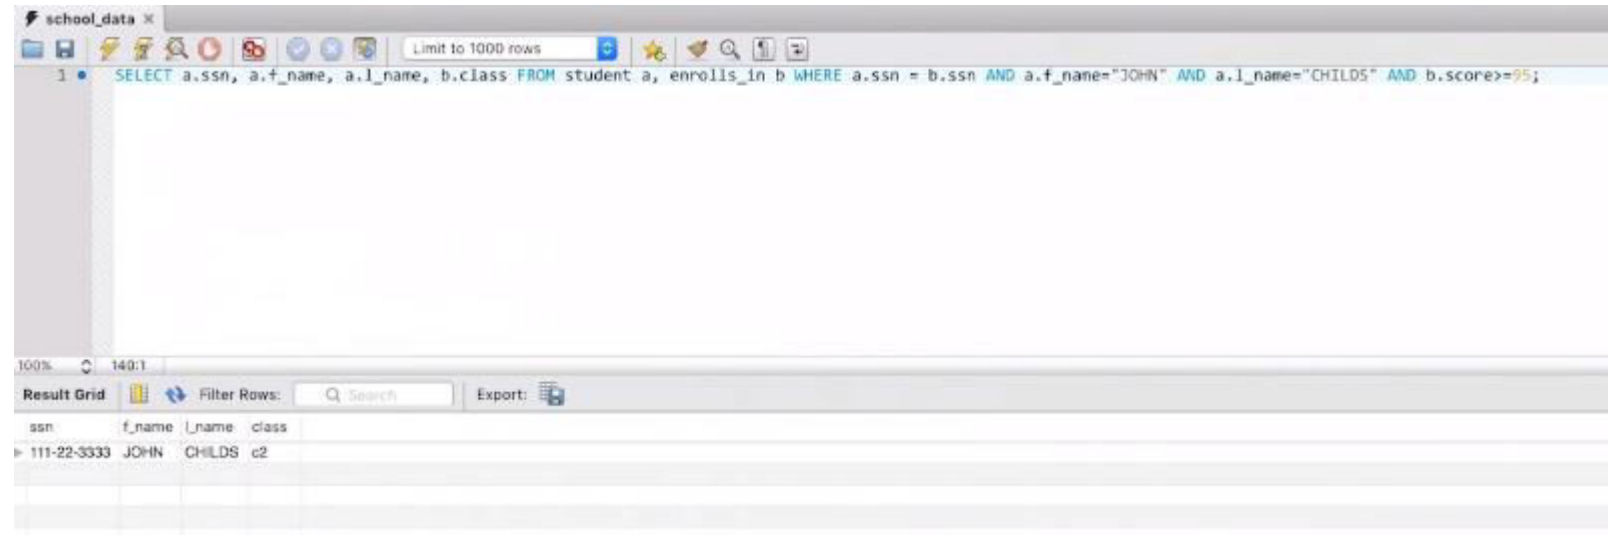
1. Get a list of student names and classes

```
SELECT f_name, l_name, class FROM Student a, Enrolls_in b  
WHERE a.ssn = b.ssn;
```

2. Get the list of classes that John Childs is enrolled in and his grade is 95 or greater

```
SELECT class FROM Student a, Enrolls_in b WHERE  
a.f_name = "JOHN" AND a.l_name = "CHILDS" AND a.ssn =  
b.ssn AND b.score >=95;
```

3. Get the names of the courses that John Childs is enrolled in
SELECT name from Student a, Enrolls_in b, Course c WHERE
a.f_name = "JOHN" AND a.l_name = "CHILDS" AND a.ssn =
b.ssn AND b.class = c.number;



The screenshot shows a database query tool window titled "school_data". The SQL query entered is: `SELECT a.ssn, a.f_name, a.l_name, b.class FROM student a, enrolls_in b WHERE a.ssn = b.ssn AND a.f_name="JOHN" AND a.l_name="CHILDS" AND b.score>=95;`. The results are displayed in a table with the following data:

ssn	f_name	l_name	class
111-22-3333	JOHN	CHILDS	c2

Working Across Multiple Tables Part 2

We can combine rows in Course with rows in Enrolls_in using an "INNER JOIN"

Join

The process of combining rows from two or more tables

Course

number	name	room
c1	Data analytics	1127
c2	Python	303
c3	corp fin	331
c4	prod. mgmt	1127
c5	Ethics	303
c6	leadership	303
c7	bus analytics	1127

Enrolls-in

ssn	class	score
111-22-3333	c1	93
123-12-1234	c1	87
111-22-3333	c2	95
222-33-4455	c3	44
555-11-7777	c1	36

rows in Enrolls-in are matched with rows in course where class has the same value as number

```
SELECT number, name, room, ssn, score FROM course
INNER JOIN Enrolls_in ON course.number = enrolls_in.class;
```

Joins can be explicit

```
SELECT number, name, room, ssn, score FROM course
INNER JOIN Enrolls_in ON course.number = enrolls_in.class;
```

or implicit

```
SELECT number, name, room, ssn, score FROM
course,enrolls_in WHERE course.number = enrolls_in.class;
```

Get the names of the courses that John is enrolled in

```
SELECT course.name FROM student
INNER JOIN enrolls_in ON student.ssn = enrolls_in.ssn
INNER JOIN course ON course.number = enrolls_in.class
WHERE f_name = "JOHN";
```

Using Python for SQL

DB API

- ➡ Set of standards for Python Database API
- ➡ Import the API module
- ➡ Acquire a connection with the database server
- ➡ Issue sql commands or call stored procedures
- ➡ Close the connection

pymysql (mysql)

- ➡ Import the API module
 - * `import pymysql`
- ➡ Acquire a connection with the database server
 - * `db = pymysql.connect("localhost","root","None")`
- ➡ Prepare a cursor object
 - * `cursor = db.cursor()`
 - * `cursor.close()`

Python and MYSQL

First import the python module containing the API

```
In [ ]: import pymysql
```

Set up a connection and create a cursor object

```
In [ ]: db = pymysql.connect("localhost","root","None",database="schooldb")
        cursor = db.cursor()
```

Execute a query and get the results

```
In [ ]: cursor.execute('show tables;')
        cursor.fetchall()
```

```
In [ ]: query = """
        SELECT course.name FROM student
        INNER JOIN enrolls_in ON student.ssn = enrolls_in.ssn
        INNER JOIN course ON course.number = enrolls_in.class
        WHERE f_name = 'JOHN';
        """
        cursor.execute(query)
        cursor.fetchall()
```

pymysql (mysql)

➡ Issue sql commands

✳ cursor.execute('show tables;')

➡ Get results

✳ cursor.fetchone()

✳ cursor.fetchall()

➡ Close the connection

✳ cursor.close()

